

Assignment 2: Building a Modeling Data Set

```
In [1]: import os
import pandas as pd
import numpy as np
%matplotlib inline
import matplotlib.pyplot as plt
import seaborn as sns
sns.set_theme()
#imports the packages used for this
```

In this assignment, you will complete the following tasks to **build a modeling dataset**:

1. Load the "adult" data set and identify the **number of rows & columns**
2. Build a new regression **label column** by winsorizing outliers
3. Replace all **missing values** with means
4. **Identify** two features with the **highest correlation with label**
5. Build appropriate **bivariate plots** between the highest correlated features and label

Part 1. Load the Data

Use the specified file name to load the data. Save it as a Pandas DataFrame called `df` .

Task: Read in the data using the `pd.read_csv()` function and save it to DataFrame `df` . Note: use the variable `filename` in your call to `pd.read_csv()` .

```
In [2]: # Do not remove or edit the line below:
filename = os.path.join(os.getcwd(), "data", "adult.data.full.asst")
```

```
In [3]: df = pd.read_csv(filename, header=0)
#loads the dataset into a data frame
```

Task: Display the shape of `df` -- that is, the number of rows and columns.

```
In [6]: df.shape
#there are 32,561 rows and 15 columns
```

```
Out[6]: (32561, 15)
```

Check your work: while we used a small subset of the `adult` dataset in the exercises, the dataset that we are using now has a substantially greater number of rows, but the same number of columns as before. You should see this reflected when you print out the dimensions of DataFrame `df` .

Task: Get a peek of the data by displaying the first few rows, as you usually do.

```
In [8]: df.head(20)
#will display the first 20 rows
```

Out[8]:

	age	workclass	fnlwgt	education	education-num	marital-status	occupation	relationship	race	sex_selfID	capital-gain	capital-loss	hours-per-week	native-country	income_binary
0	39.0	State-gov	77516	Bachelors	13	Never-married	Adm-clerical	Not-in-family	White	Non-Female	2174	0	40.0	United-States	<=50K
1	50.0	Self-emp-not-inc	83311	Bachelors	13	Married-civ-spouse	Exec-managerial	Husband	White	Non-Female	0	0	13.0	United-States	<=50K
2	38.0	Private	215646	HS-grad	9	Divorced	Handlers-cleaners	Not-in-family	White	Non-Female	0	0	40.0	United-States	<=50K
3	53.0	Private	234721	11th	7	Married-civ-spouse	Handlers-cleaners	Husband	Black	Non-Female	0	0	40.0	United-States	<=50K
4	28.0	Private	338409	Bachelors	13	Married-civ-spouse	Prof-specialty	Wife	Black	Female	0	0	40.0	Cuba	<=50K
5	37.0	Private	284582	Masters	14	Married-civ-spouse	Exec-managerial	Wife	White	Female	0	0	40.0	United-States	<=50K
6	49.0	Private	160187	9th	5	Married-spouse-absent	Other-service	Not-in-family	Black	Female	0	0	16.0	Jamaica	<=50K
7	52.0	Self-emp-not-inc	209642	HS-grad	9	Married-civ-spouse	Exec-managerial	Husband	White	Non-Female	0	0	45.0	United-States	>50K
8	31.0	Private	45781	Masters	14	Never-married	Prof-specialty	Not-in-family	White	Female	14084	0	50.0	United-States	>50K
9	42.0	Private	159449	Bachelors	13	Married-civ-spouse	Exec-managerial	Husband	White	Non-Female	5178	0	40.0	United-States	>50K
10	37.0	Private	280464	Some-college	10	Married-civ-spouse	Exec-managerial	Husband	Black	Non-Female	0	0	80.0	United-States	>50K
11	30.0	State-gov	141297	Bachelors	13	Married-civ-spouse	Prof-specialty	Husband	Asian-Pac-Islander	Non-Female	0	0	40.0	India	>50K
12	23.0	Private	122272	Bachelors	13	Never-married	Adm-clerical	Own-child	White	Female	0	0	30.0	United-States	<=50K
13	32.0	Private	205019	Assoc-acdm	12	Never-married	Sales	Not-in-family	Black	Non-Female	0	0	50.0	United-States	<=50K
14	40.0	Private	121772	Assoc-voc	11	Married-civ-spouse	Craft-repair	Husband	Asian-Pac-Islander	Non-Female	0	0	40.0	NaN	>50K
15	34.0	Private	245487	7th-8th	4	Married-civ-spouse	Transport-moving	Husband	Amer-Indian-Inuit	Non-Female	0	0	45.0	Mexico	<=50K
16	25.0	Self-emp-not-inc	176756	HS-grad	9	Never-married	Farming-fishing	Own-child	White	Non-Female	0	0	35.0	United-States	<=50K
17	32.0	Private	186824	HS-grad	9	Never-married	Machine-op-inspct	Unmarried	White	Non-Female	0	0	NaN	United-States	<=50K
18	38.0	Private	28887	11th	7	Married-civ-spouse	Sales	Husband	White	Non-Female	0	0	50.0	United-States	<=50K
19	43.0	Self-emp-not-inc	292175	Masters	14	Divorced	Exec-managerial	Unmarried	White	Female	0	0	45.0	United-States	>50K

Part 2. Create a (Winsorized) Label Column

Assume that your goal is to use this dataset to fit a regression model that predicts the number of years of education that a person has had.

We'd like to create a new version of the `education-num` column, in which we replace the outlier values of `education-num` (on both sides of the range -- the low end as well as the high end). We will replace the outliers with the corresponding percentile value, as we did in the exercises. That is, if we wish to replace any value below, say, the 1.234-th percentile, we shall replace all such (various) values by the exact same value in our data -- the value such that 1.234% of data lies below it.

We will need to import the `stats` module from the `scipy` package:

```
In [9]: import scipy.stats as stats
```

Task: Create a new column, titled `label`, by winsorizing the `education-num` column with the top and bottom 1% percentile values.

```
In [11]: # a new column label that uses the winsorize method
df['label'] = stats.mstats.winsorize(df['education-num'], limits=[0.01, 0.01])
```

Let's verify that a new column got added to the DataFrame:

```
In [12]: df.head()
```

```
Out[12]:
```

	age	workclass	fnlwgt	education	education-num	marital-status	occupation	relationship	race	sex_selfID	capital-gain	capital-loss	hours-per-week	native-country	income_binary	label
0	39.0	State-gov	77516	Bachelors	13	Never-married	Adm-clerical	Not-in-family	White	Non-Female	2174	0	40.0	United-States	<=50K	13
1	50.0	Self-emp-not-inc	83311	Bachelors	13	Married-civ-spouse	Exec-managerial	Husband	White	Non-Female	0	0	13.0	United-States	<=50K	13
2	38.0	Private	215646	HS-grad	9	Divorced	Handlers-cleaners	Not-in-family	White	Non-Female	0	0	40.0	United-States	<=50K	9
3	53.0	Private	234721	11th	7	Married-civ-spouse	Handlers-cleaners	Husband	Black	Non-Female	0	0	40.0	United-States	<=50K	7
4	28.0	Private	338409	Bachelors	13	Married-civ-spouse	Prof-specialty	Wife	Black	Female	0	0	40.0	Cuba	<=50K	13

An interesting thing to think about: take a look at the data and notice that for the first five rows, the values of the `education-num` column and its winsorized version -- `label` -- are the same. Does this mean that winsorization did not work? Or are there discrepancies further down the list of rows, where we cannot see them?

Task: Check that the values of `education-num` and `label` are *not* identical. You may do this by subtracting the two columns and then listing the unique values of the result. If you see values other than zero, it means *some* change did happen, as we would expect.

```
In [29]: check = df['education-num'] - df['label']
check.unique()
check
```

```
Out[29]: 0
1
2
3
4
..
32556 0
32557 0
32558 0
32559 0
32560 0
Length: 32561, dtype: int64
```

Part 3. Replace the Missing Values With Means

a. Identifying missingness

Task: Check if a given value in any data cell is missing, and sum up the resulting values (`True` / `False`) by columns. Assign the results to variable `nan_count` . Print the results.

```
In [30]: df.isnull().values.any()
nan_count = np.sum(df.isnull(), axis = 0)
nan_count
```

```
Out[30]: age                162
workclass            1836
fnlwgt                0
education             0
education-num         0
marital-status        0
occupation           1843
relationship          0
race                  0
sex_selfID            0
capital-gain          0
capital-loss          0
hours-per-week       325
native-country        583
income_binary         0
label                 0
dtype: int64
```

Replacing the missing values with the mean only makes sense for the numerically valued columns (and not for strings). Hence, we will focus on the `age` and `hours-per-week` columns.

b. Keeping record of the missingness: creating dummy variables

As a first step, you will now create dummy variables indicating missingness of the values.

Task: Store the `True` / `False` series that indicate missingness of any value in `age` as a new column called `age_na` . Store the `True` / `False` series that indicate missingness of every value of `hours-per-week` as a new column called `hours-per-week_na` .

```
In [31]: # isnull will count the missing values
df['age_na'] = df['age'].isnull()
df['hours-per-week_na'] = df['hours-per-week'].isnull()
df.head()
```

Out[31]:

	age	workclass	fnlwgt	education	education-num	marital-status	occupation	relationship	race	sex	selfid	capital-gain	capital-loss	hours-per-week	native-country	income_binary	label	age_na	hours per week_na
0	39.0	State-gov	77516	Bachelors	13	Never-married	Adm-clerical	Not-in-family	White	Non-Female		2174	0	40.0	United-States	<=50K	13	False	False
1	50.0	Self-emp-not-inc	83311	Bachelors	13	Married-civ-spouse	Exec-managerial	Husband	White	Non-Female		0	0	13.0	United-States	<=50K	13	False	False
2	38.0	Private	215646	HS-grad	9	Divorced	Handlers-cleaners	Not-in-family	White	Non-Female		0	0	40.0	United-States	<=50K	9	False	False
3	53.0	Private	234721	11th	7	Married-civ-spouse	Handlers-cleaners	Husband	Black	Non-Female		0	0	40.0	United-States	<=50K	7	False	False
4	28.0	Private	338409	Bachelors	13	Married-civ-spouse	Prof-specialty	Wife	Black	Female		0	0	40.0	Cuba	<=50K	13	False	False

c. Replacing the missing values with mean values of the column

Task: Fill the missing values of the `age` and `hours-per-week` columns with the mean value of the corresponding column.

```
In [34]: # compute mean for all non null age values
mean_ages=df['age'].mean()
print("mean value for all age columns: " + str(mean_ages))

# fill all missing values with the mean
df['age'].fillna(value=mean_ages, inplace=True)

# compute mean for all non null hpw values
mean_hours=df['hours-per-week'].mean()
print("mean value for all hours-per-week columns: " + str(mean_hours))

# fill all missing values with the mean
df['hours-per-week'].fillna(value=mean_hours, inplace=True)
```

mean value for all age columns: 38.58921571653446
mean value for all hours-per-week columns: 40.450428092815486

Ungraded Task: Check your results. Display the sum of missing values for the `age` column (or reuse the code for listing total numbers of all missing values that you wrote before, if it worked).

```
In [ ]: # YOUR CODE HERE - this cell will not be graded
```

Part 4. Identify Features With the Highest Correlation With the Label

Your next goal is to figure out which features in the data correlate most with the label.

In the next few cells, we will demonstrate how to use Pandas `corr()` method to get a list of correlation coefficients between the `label` and all other (numerical) features. To learn more about the `corr()` method, consult the online [documentation \(https://pandas.pydata.org/docs/reference/api/pandas.DataFrame.corr.html\)](https://pandas.pydata.org/docs/reference/api/pandas.DataFrame.corr.html).

Let's first glance at what the `corr()` method does:

```
In [35]: df.corr()
```

Out[35]:

	age	fnlwgt	education-num	capital-gain	capital-loss	hours-per-week	label	age_na	hours-per-week_na
age	1.000000e+00	-0.076085	0.036685	0.124705	0.057478	6.657191e-02	0.038549	7.101579e-18	-4.325250e-05
fnlwgt	-7.608468e-02	1.000000	-0.043195	-0.002234	-0.010252	-1.804716e-02	-0.042134	-9.015193e-03	-5.6769619e-03
education-num	3.668517e-02	-0.043195	1.000000	0.167089	0.079923	1.465533e-01	0.999182	-1.708530e-03	-5.670679e-03
capital-gain	1.247046e-01	-0.002234	0.167089	1.000000	-0.055138	1.009947e-01	0.168202	-5.313515e-03	4.981172e-03
capital-loss	5.747841e-02	-0.010252	0.079923	-0.055138	1.000000	5.420158e-02	0.080453	-7.205893e-03	-1.511760e-03
hours-per-week	6.657191e-02	-0.018047	0.146553	0.100995	0.054202	1.000000e+00	0.147275	2.254277e-03	7.385613e-17
label	3.854869e-02	-0.042134	0.999182	0.168202	0.080453	1.472753e-01	1.000000	-1.955584e-03	-5.811006e-03
age_na	7.101579e-18	-0.009015	-0.001709	-0.005314	-0.007206	2.254277e-03	-0.001956	1.000000e+00	-2.709086e-03
hours-per-week_na	-4.325250e-05	-0.005770	-0.005671	0.004981	-0.001512	7.385613e-17	-0.005811	-2.709086e-03	1.000000e+00

The result is a computed *correlation matrix*. The values on the diagonal are all equal to 1, and the matrix is symmetrical with respect to the diagonal.

We only need to observe correlations of all features with the column `label` (as opposed to every possible pairwise correlation). So let's query the `label` column of this matrix:

```
In [36]: df.corr()['label']

Out[36]: age                0.038549
fnlwgt             -0.042134
education-num       0.999182
capital-gain        0.168202
capital-loss        0.080453
hours-per-week      0.147275
label              1.000000
age_na             -0.001956
hours-per-week_na  -0.005811
Name: label, dtype: float64
```

This is good, but contains two values too many: we do not need to observe the correlation of `label` with itself, and moreover we are not interested in the correlation between the label and `education-num` (recall that `label` is a winsorized version of the `education-num`). So we will exclude these two values using the Pandas `drop()` method:

```
In [37]: exclude = ['label', 'education-num']
df.corr()['label'].drop(exclude, axis = 0)

Out[37]: age                0.038549
fnlwgt             -0.042134
capital-gain        0.168202
capital-loss        0.080453
hours-per-week      0.147275
age_na             -0.001956
hours-per-week_na  -0.005811
Name: label, dtype: float64
```

Task: The code below performs the same operation above, but saves the result to variable `corrs`. Sort the values in `corrs` in descending order. Use the Pandas method `sort_values()` to accomplish this task. For more information on how to use the `sort_values()` method, consult the online [documentation](https://pandas.pydata.org/docs/reference/api/pandas.Series.sort_values.html) (https://pandas.pydata.org/docs/reference/api/pandas.Series.sort_values.html).

```
In [44]: # Do not remove or edit the line below:
corrs = df.corr()['label'].drop(exclude, axis = 0)

corrs_sorted = corrs.sort_values(ascending=False)

Out[44]: age                0.038549
fnlwgt             -0.042134
capital-gain        0.168202
capital-loss        0.080453
hours-per-week      0.147275
age_na             -0.001956
hours-per-week_na  -0.005811
Name: label, dtype: float64
```

Task: Save the *column names* for the top-2 correlation values into a Python list called `top_two_corr`

Tip: `corrs_sorted` is a Pandas `Series` object, in which column names are the *index*. Once you find the column names, use the Python `list()` method to convert the values into a Python `list`.

```
In [55]: top_two_corr = list(corrs_sorted)
top_two_corr

Out[55]: [0.1682023900104458,
0.1472752942611272,
0.0804534806659499,
0.03854869188977444,
-0.001955583625737097,
-0.005811006138601126,
-0.04213404090183742]
```

Part 5. Produce Bivariate Plots for the Label and Its Top Correlates

We will use the `pairplot()` function in `seaborn` to plot the relationships between the two features we identified and the label.

Task: Create a `DataFrame` named `df_sub` that contains only these three columns from `DataFrame df`: the label, and the two columns which correlate with it the most.

Tip: You can use the variable `top_two_corrs` in your solution.

```
In [63]: df_sub = df[top_two_corrs].copy()
```

```
-----
KeyError                                Traceback (most recent call last)
<ipython-input-63-d30efeee63b7> in <module>()
----> 1 df_sub = df[top_two_corrs].copy()

~/.local/lib/python3.6/site-packages/pandas/core/frame.py in __getitem__(self, key)
 2910     if is_iterator(key):
 2911         key = list(key)
-> 2912     indexer = self.loc._get_listlike_indexer(key, axis=1, raise_missing=True)[1]
 2913
 2914     # take() does not accept boolean indexers

~/.local/lib/python3.6/site-packages/pandas/core/indexing.py in _get_listlike_indexer(self, key, axis, raise_missing)
 1252     keyarr, indexer, new_indexer = ax._reindex_non_unique(keyarr)
 1253
-> 1254     self._validate_read_indexer(keyarr, indexer, axis, raise_missing=raise_missing)
 1255     return keyarr, indexer
 1256

~/.local/lib/python3.6/site-packages/pandas/core/indexing.py in _validate_read_indexer(self, key, indexer, axis, raise_missing)
 1296     if missing == len(indexer):
 1297         axis_name = self.obj._get_axis_name(axis)
-> 1298         raise KeyError(f"None of [{key}] are in the [{axis_name}]")
 1299
 1300     # We (temporarily) allow for some missing keys with .loc, except in

KeyError: "None of [Float64Index([ 0.1682023900104458,    0.1472752942611272,\n188977444,\n                                -0.001955583625737097, -0.005811006138601126,\n                                0.0804534806659499,    0.03854869\n                                -0.04213404090183742],\n                                dtype\n='float64')] are in the [columns]"
```

Task: Create a seaborn pairplot of the data subset you just created.

```
In [64]: sns.pairplot(data=df_sub)
```

```
-----
NameError                                Traceback (most recent call last)
<ipython-input-64-b05a3754fbcd> in <module>()
----> 1 sns.pairplot(data=df_sub)

NameError: name 'df_sub' is not defined
```

This one is not very easy to make sense of: the points overlap, but we do not have visibility into how densely they are stacked together.

Task: Repeat the `pairplot` exercise, this time specifying the *kernel density estimator* as the *kind* of the plot.

Tip: Use `kind = 'kde'` as a parameter of the `pairplot()` function. You could also specify `corner=True` to make sure you don't plot redundant (symmetrical) plots.

Note: This will take a while to run and produce a plot.

```
In [ ]: # YOUR CODE HERE
```

Think about the possible interpretations of these plots. (Recall that our label encodes `education`, in number of years).

Here is an example of the kind of stories this data seems to be telling. It appears as though hours per week are stacked around the typical 40-hour value, and that this value of weekly hours dominates regardless of the level of education. However, it seems that it is somewhat less typical for people with lower levels of formal education to be working over 65 hours a week.

Analysis: Try to interpret what you see in this plot, as well as the one depicting the relationship between 'capital gain' and the levels of education, and see what kind of patterns you are noticing. Moreover, is there something odd that raises red flags and makes you think the data or our handling of it may be invalid? Is there something that, on the contrary, satisfies your intuition, thereby providing a 'sanity check'? These are the kind of questions that are useful to ask yourself as you are looking at the data distributions and pairwise relationships. Record your findings in the cell below.

<Double click this Markdown cell to make it editable, and record your findings here.>