

Break Through Tech Machine Learning Week 3

April 14, 2023

1 Assignment 3: Building a Decision Tree After Feature Transformations

```
[1]: import pandas as pd
import numpy as np
import os
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import OneHotEncoder
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import accuracy_score
```

In this assignment, you will implement the following steps to build a Decision Tree classification model:

1. Load the "cell2celltrain" data set
2. Convert categorical features to one-hot encoded values
3. Split the data into training and test sets
4. Fit a Decision Tree classifier and evaluate the accuracy of its predictions
5. Plot the training set accuracy

1.1 Part 1. Load the Data Set

We will work with the "cell2celltrain" data set.

```
[2]: # Do not remove or edit the line below:
filename = os.path.join(os.getcwd(), "data", "cell2celltrain.csv")
```

Task: Load the data and save it to DataFrame df.

```
[3]: # YOUR CODE HERE

# Solution
df = pd.read_csv(filename, header=0)
```

Task: Display the shape of df -- that is, the number of records (rows) and variables (columns)

```
[4]: # YOUR CODE HERE
```

```
# Solution
df.shape
```

[4]: (51047, 58)

For the purpose of this assignment, we will remove the Married column due to missing values

[5]: `df.drop(columns = ['Married'], inplace=True)`

1.2 Part 2. One-Hot Encode Categorical Values

To implement a decision tree model, we must first transform the string-valued categorical features into numerical boolean values using one-hot encoding.

1.2.1 a. Find the Columns Containing String Values

[6]: `df.dtypes`

[6]:

CustomerID	int64
Churn	bool
ServiceArea	object
ChildrenInHH	bool
HandsetRefurbished	bool
HandsetWebCapable	bool
TruckOwner	bool
RVOwner	bool
HomeownershipKnown	bool
BuysViaMailOrder	bool
RespondsToMailOffers	bool
OptOutMailings	bool
NonUSTravel	bool
OwnsComputer	bool
HasCreditCard	bool
NewCellphoneUser	bool
NotNewCellphoneUser	bool
OwnsMotorcycle	bool
MadeCallToRetentionTeam	bool
CreditRating	object
PrizmCode	object
Occupation	object
MonthlyRevenue	float64
MonthlyMinutes	float64
TotalRecurringCharge	float64
DirectorAssistedCalls	float64
OverageMinutes	float64
RoamingCalls	float64
PercChangeMinutes	float64
PercChangeRevenues	float64
DroppedCalls	float64

BlockedCalls	float64
UnansweredCalls	float64
CustomerCareCalls	float64
ThreewayCalls	float64
ReceivedCalls	float64
OutboundCalls	float64
InboundCalls	float64
PeakCallsInOut	float64
OffPeakCallsInOut	float64
DroppedBlockedCalls	float64
CallForwardingCalls	float64
CallWaitingCalls	float64
MonthsInService	float64
UniqueSubs	float64
ActiveSubs	float64
Handsets	float64
HandsetModels	float64
CurrentEquipmentDays	float64
AgeHH1	float64
AgeHH2	float64
RetentionCalls	float64
RetentionOffersAccepted	float64
ReferralsMadeBySubscriber	float64
IncomeGroup	float64
AdjustmentsToCreditRating	float64
HandsetPrice	float64
dtype:	object

Task: Add all of the column names whos values are of type 'object' to a list named `to_encode`.

```
[7]: # YOUR CODE HERE

#solution:
to_encode = list(df.select_dtypes(include=['object']).columns)
```

Let's take a closer look at the candidates for one-hot encoding:

```
[8]: df[to_encode].nunique()

ServiceArea      747
CreditRating      7
PrizmCode         4
Occupation        8
dtype: int64
```

For all of the columns except for `ServiceArea`, it should be straightforward to replace a given column with a set of several new binary columns for each unique value. However, let's first deal with the special case of `ServiceArea`.

1.2.2 b. One Hot-Encoding 'ServiceArea': The Top 10 Values

Take a look at the number of unique values of the ServiceArea column. There are too many unique values in the ServiceArea column to attempt to create a new binary indicator column per value! One thing we could do is to see if some of the values in ServiceArea are occurring frequently. We will then one-hot encode just those frequent values.

Task: Get the top 10 most frequent values in 'ServiceArea' and store them in list top_10_SA.

```
[9]: # YOUR CODE HERE
```

```
### Solution:
top_10_SA = list(df['ServiceArea'].value_counts().head(10).index)
top_10_SA
```

```
[9]: ['NYCBRO917',
      'HOUHOU281',
      'DALDAL214',
      'NYCMAN917',
      'APCFCH703',
      'DALFTW817',
      'SANSAN210',
      'APCSIL301',
      'SANAUS512',
      'SFROAK510']
```

Task: Write a for loop that loops through every value in top_10_SA and creates one-hot encoded columns, titled 'ServiceArea + '_' + <service area value>'. For example, there will be a column named 'ServiceArea_NYCBRO917'. Use the NumPy np.where() function to accomplish this.

```
[10]: # YOUR CODE HERE
```

```
# Solution
for value in top_10_SA:

    ## Create columns and their values
    df['ServiceArea_' + value] = np.where(df['ServiceArea']==value,1,0)
```

Task: 1. Drop the original, multi-valued ServiceArea column from the DataFrame df. 2. Remove 'ServiceArea' from the to_encode list.

```
[11]: # YOUR CODE HERE
```

```
# Remove the original column from your DataFrame df
df.drop(columns = 'ServiceArea', inplace=True)

# Remove from list to_encode
to_encode.remove('ServiceArea')
```

```
[12]: df.head()
```

[12]:

	CustomerID	Churn	ChildrenInHH	HandsetRefurbished	HandsetWebCapable	\
0	3000002	True	False	False	True	
1	3000010	True	True	False	False	
2	3000014	False	True	False	False	
3	3000022	False	False	False	True	
4	3000026	True	False	False	False	

	TruckOwner	RVOwner	HomeownershipKnown	BuysViaMailOrder	\
0	False	False	True	True	
1	False	False	True	True	
2	False	False	False	False	
3	False	False	True	True	
4	False	False	True	True	

	RespondsToMailOffers	...	ServiceArea_NYCBRO917	ServiceArea_HOUHOU281	\
0	True	...	0	0	
1	True	...	0	0	
2	False	...	0	0	
3	True	...	0	0	
4	True	...	0	0	

	ServiceArea_DALDAL214	ServiceArea_NYCMAN917	ServiceArea_APCFCH703	\
0	0	0	0	
1	0	0	0	
2	0	0	0	
3	0	0	0	
4	0	0	0	

	ServiceArea_DALFTW817	ServiceArea_SANSAN210	ServiceArea_APCSIL301	\
0	0	0	0	
1	0	0	0	
2	0	0	0	
3	0	0	0	
4	0	0	0	

	ServiceArea_SANAUS512	ServiceArea_SFROAK510
0	0	0
1	0	0
2	0	0
3	0	0
4	0	0

[5 rows x 66 columns]

1.2.3 c. One Hot-Encoding all Remaining Columns: All Unique Values per Column

All other columns in `to_encode` have reasonably small numbers of unique values, so we are going to simply one-hot encode every unique value of those columns.

Task: In the code cell below, iterate over column names and create new columns for all unique values. 1. Use a loop to loop over the column names in `to_encode` 2. In the loop: 1. Use the Pandas `pd.get_dummies()` function and save the result to variable `temp_df` 2. Use `df.join` to join `temp_df` with DataFrame `df`

[13]: `# YOUR CODE HERE`

```
# SOLUTION
for colname in to_encode:
    temp_df = pd.get_dummies(df[colname], prefix=colname + '_')
    df = df.join(temp_df)
```

[14]: `df.head()`

```
[14]: CustomerID  Churn  ChildrenInHH  HandsetRefurbished  HandsetWebCapable  \
0      3000002   True           False                False                True
1      3000010   True           True                 False                False
2      3000014  False           True                 False                False
3      3000022  False           False                False                True
4      3000026   True           False                False                False

      TruckOwner  RVOwner  HomeownershipKnown  BuysViaMailOrder  \
0         False   False                True                True
1         False   False                True                True
2         False   False               False               False
3         False   False                True                True
4         False   False                True                True

      RespondsToMailOffers  ...  PrizmCode__Suburban  PrizmCode__Town  \
0                True  ...                1                0
1                True  ...                1                0
2               False  ...                0                1
3                True  ...                0                0
4                True  ...                0                0

      Occupation__Clerical  Occupation__Crafts  Occupation__Homemaker  \
0                        0                    0                    0
1                        0                    0                    0
2                        0                    1                    0
3                        0                    0                    0
4                        0                    0                    0

      Occupation__Other  Occupation__Professional  Occupation__Retired  \
0                        0                        1                    0
1                        0                        1                    0
```

2	0	0	0
3	1	0	0
4	0	1	0

	Occupation__Self	Occupation__Student
0	0	0
1	0	0
2	0	0
3	0	0
4	0	0

[5 rows x 85 columns]

Task: Remove all the original columns from DataFrame df

[15]: *# YOUR CODE HERE*

Solution

df.drop(columns = to_encode ,axis=1, inplace=True)

[16]: df.columns

[16]: Index(['CustomerID', 'Churn', 'ChildrenInHH', 'HandsetRefurbished',
 'HandsetWebCapable', 'TruckOwner', 'RVOwner', 'HomeownershipKnown',
 'BuysViaMailOrder', 'RespondsToMailOffers', 'OptOutMailings',
 'NonUSTravel', 'OwnsComputer', 'HasCreditCard', 'NewCellphoneUser',
 'NotNewCellphoneUser', 'OwnsMotorcycle', 'MadeCallToRetentionTeam',
 'MonthlyRevenue', 'MonthlyMinutes', 'TotalRecurringCharge',
 'DirectorAssistedCalls', 'OverageMinutes', 'RoamingCalls',
 'PercChangeMinutes', 'PercChangeRevenues', 'DroppedCalls',
 'BlockedCalls', 'UnansweredCalls', 'CustomerCareCalls', 'ThreewayCalls',
 'ReceivedCalls', 'OutboundCalls', 'InboundCalls', 'PeakCallsInOut',
 'OffPeakCallsInOut', 'DroppedBlockedCalls', 'CallForwardingCalls',
 'CallWaitingCalls', 'MonthsInService', 'UniqueSubs', 'ActiveSubs',
 'Handsets', 'HandsetModels', 'CurrentEquipmentDays', 'AgeHH1', 'AgeHH2',
 'RetentionCalls', 'RetentionOffersAccepted',
 'ReferralsMadeBySubscriber', 'IncomeGroup', 'AdjustmentsToCreditRating',
 'HandsetPrice', 'ServiceArea_NYCBRO917', 'ServiceArea_HOUHOU281',
 'ServiceArea_DALDAL214', 'ServiceArea_NYCMAN917',
 'ServiceArea_APCFCH703', 'ServiceArea_DALFTW817',
 'ServiceArea_SANSAN210', 'ServiceArea_APCSIL301',
 'ServiceArea_SANAUS512', 'ServiceArea_SFROAK510',
 'CreditRating__1-Highest', 'CreditRating__2-High',
 'CreditRating__3-Good', 'CreditRating__4-Medium', 'CreditRating__5-Low',
 'CreditRating__6-VeryLow', 'CreditRating__7-Lowest', 'PrizmCode__Other',
 'PrizmCode__Rural', 'PrizmCode__Suburban', 'PrizmCode__Town',
 'Occupation__Clerical', 'Occupation__Crafts', 'Occupation__Homemaker',
 'Occupation__Other', 'Occupation__Professional', 'Occupation__Retired',
 'Occupation__Self', 'Occupation__Student'],

```
dtype='object')
```

Check that the data does not contain any missing values. The absence of missing values is necessary for training a Decision Tree model.

```
[17]: # YOUR CODE HERE
```

```
# solution
df.isnull().values.any()
```

```
[17]: False
```

1.3 Part 3: Create Labeled Examples from the Data Set

Task: Create labeled examples from DataFrame df. In the code cell below carry out the following steps:

- Get the Churn column from DataFrame df and assign it to the variable y. This will be our label.
- Get all other columns from DataFrame df and assign them to the variable X. These will be our features.

```
[18]: # YOUR CODE HERE
```

```
### Solution:
y = df['Churn']
X = df.drop(columns = 'Churn', axis=1)
```

1.4 Part 4: Create Training and Test Data Sets

Task: In the code cell below create training and test sets out of the labeled examples.

1. Use Scikit-learn's `train_test_split()` function to create the data sets.
2. Specify:
 - A test set that is 30 percent (.30) of the size of the data set.
 - A seed value of '123'.

```
[19]: # YOUR CODE HERE
```

```
### Solution:
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3,
→random_state=123)
```

Check that the dimensions of the training and test datasets are what you expected:

```
[20]: print(X_train.shape)
print(X_test.shape)
```

```
(35732, 81)
(15315, 81)
```


1.5 Part 5. Fit a Decision Tree Classifier and Evaluate the Model

The code cell below contains a shell of a function named `train_test_DT()`. This function should train a Decision Tree classifier on the training data, test the resulting model on the test data, and compute and return the accuracy score of the resulting predicted class labels on the test data.

Task: Complete the function to make it work.

```
[21]: def train_test_DT(X_train, X_test, y_train, y_test, leaf, depth,
    crit='entropy'):
    """
    Fit a Decision Tree classifier to the training data X_train, y_train.
    Return the accuracy of resulting predictions on the test set.
    Parameters:
        leaf := The minimum number of samples required to be at a leaf node
        depth := The maximum depth of the tree
        crit := The function to be used to measure the quality of a split.
    →Default: gini.
    """

    # 1. Create the Scikit-learn DecisionTreeClassifier model object below
    →and assign to variable 'model'
    # YOUR CODE HERE

    ### SOLUTION
    model = DecisionTreeClassifier(criterion = crit, max_depth = depth,
    →min_samples_leaf = leaf)

    # 2. Fit the model to the training data below
    # YOUR CODE HERE

    # SOLUTION
    model.fit(X_train, y_train)

    # 3. Make predictions on the test data and assign the result to the
    →variable 'class_label_predictions' below
    # YOUR CODE HERE

    #SOLUTION
    class_label_predictions = model.predict(X_test)

    # 4. Compute the accuracy and save the result to the variable 'acc_score'
    →below
    # YOUR CODE HERE

    #SOLUTION
    acc_score = accuracy_score(y_test, class_label_predictions)
```

```
return acc_score
```

1.5.1 Train on Different Hyperparameter Values

Task: Train two Decision Tree classifiers using your function.

- one with a low value of depth
- one high value of depth

Specify the minimum number of samples at the leaf node to be equal to 1 for both trees.
Save the resulting accuracy scores to list acc. Print the list.

```
[22]: depth1= 2 # YOUR CODE HERE (solutions will vary)
      depth2 = 5 # YOUR CODE HERE (solutions will vary)
      leaf = 1

      max_depth_range = [depth1, depth2]
      acc = []

      # YOUR CODE HERE

      ### Solution:

      for md in max_depth_range:
          score = train_test_DT(X_train, X_test, y_train, y_test, 1, md)
          print('Max Depth=' + str(md) + ', accuracy score: ' + str(score))
          acc.append(float(score))
```

Max Depth=2, accuracy score: 0.7107411034933072


Max Depth=5, accuracy score: 0.715834149526608

Task: Visualize the results (Hint: use a seaborn lineplot).

```
[23]: fig = plt.figure()
      ax = fig.add_subplot(111)

      # YOUR CODE HERE
      # solution
      p = sns.lineplot(x=max_depth_range, y=acc, marker='o', label = 'Full training_
      ↳set')

      plt.title('Test set accuracy of the DT predictions, for $max\_depth\in\{8,
      ↳32\}$')
      ax.set_xlabel('max_depth')
      ax.set_ylabel('Accuracy')
      plt.show()
```



Analysis: Experiment with different values for max_depth . Add these values to the list max_depth_range (i.e. change the values, create a list containing more values), retrain your model and rerun with the visualization cell above. Compare the different accuracy scores.

Once you find the best value for max_depth , experiment with different values for leaf and compare the different accuracy scores.

Is there one model configuration that yields the best score? Record your findings in the cell below.