

## Lab 8: Implement Your Machine Learning Project Plan

In this lab assignment, you will implement the machine learning project plan you created in the written assignment. You will:

1. Load your data set and save it to a Pandas DataFrame.
2. Perform exploratory data analysis on your data to determine which feature engineering and data preparation techniques you will use.
3. Prepare your data for your model and create features and a label.
4. Fit your model to the training data and evaluate your model.
5. Improve your model by performing model selection and/or feature selection techniques to find best model for your problem.

### Import Packages

Before you get started, import a few packages.

```
In [1]: import pandas as pd
import numpy as np
import os
import matplotlib.pyplot as plt
import seaborn as sns
```

**Task:** In the code cell below, import additional packages that you have used in this course that you will need for this task.

```
In [2]: # these are the additional packages that will ne needed for this task
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import plot_roc_curve, accuracy_score, roc_auc_score
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.pipeline import Pipeline
```

### Part 1: Load the Data Set

You have chosen to work with one of four data sets. The data sets are located in a folder named "data." The file names of the three data sets are as follows:

- The "adult" data set that contains Census information from 1994 is located in file `adultData.csv`
- The airbnb NYC "listings" data set is located in file `airbnbListingsData.csv`
- The World Happiness Report (WHR) data set is located in file `WHR2018Chapter2OnlineData.csv`
- The book review data set is located in file `bookReviewsData.csv`

**Task:** In the code cell below, use the same method you have been using to load your data using `pd.read_csv()` and save it to DataFrame `df`.

```
In [3]: # File names of the book review data set
adultDataSet_filename = os.path.join(os.getcwd(), "data", "adultData.csv")

df = pd.read_csv(adultDataSet_filename, header=0)

df.head()
```

Out[3]:

	age	workclass	fnlwgt	education	education-num	marital-status	occupation	relationship	race	sex_selfID	capital-gain	capital-loss	hours-per-week	native-country	income_binary
0	39.0	State-gov	77516	Bachelors	13	Never-married	Adm-clerical	Not-in-family	White	Non-Female	2174	0	40.0	United-States	<=50K
1	50.0	Self-emp-not-inc	83311	Bachelors	13	Married-civ-spouse	Exec-managerial	Husband	White	Non-Female	0	0	13.0	United-States	<=50K
2	38.0	Private	215646	HS-grad	9	Divorced	Handlers-cleaners	Not-in-family	White	Non-Female	0	0	40.0	United-States	<=50K
3	53.0	Private	234721	11th	7	Married-civ-spouse	Handlers-cleaners	Husband	Black	Non-Female	0	0	40.0	United-States	<=50K
4	28.0	Private	338409	Bachelors	13	Married-civ-spouse	Prof-specialty	Wife	Black	Female	0	0	40.0	Cuba	<=50K

## Part 2: Exploratory Data Analysis

The next step is to inspect and analyze your data set with your machine learning problem and project plan in mind.

This step will help you determine data preparation and feature engineering techniques you will need to apply to your data to build a balanced modeling data set for your problem and model. These data preparation techniques may include:

- addressing missingness, such as replacing missing values with means
- renaming features and labels
- finding and replacing outliers
- performing winsorization if needed
- performing one-hot encoding on categorical features
- performing vectorization for an NLP problem
- addressing class imbalance in your data sample to promote fair AI

Think of the different techniques you have used to inspect and analyze your data in this course. These include using Pandas to apply data filters, using the Pandas `describe()` method to get insight into key statistics for each column, using the Pandas `dtypes` property to inspect the data type of each column, and using Matplotlib and Seaborn to detect outliers and visualize relationships between features and labels. If you are working on a classification problem, use techniques you have learned to determine if there is class imbalance.

**Task:** Use the techniques you have learned in this course to inspect and analyze your data.

**Note:** You can add code cells if needed by going to the **Insert** menu and clicking on **Insert Cell Below** in the drop-down menu.

```
In [4]: summary = df.describe()
print(summary)
```

	age	fnlwgt	education-num	capital-gain	capital-loss	\
count	32399.000000	3.256100e+04	32561.000000	32561.000000	32561.000000	
mean	38.589216	1.897784e+05	10.080679	615.907773	87.303830	
std	13.647862	1.055500e+05	2.572720	2420.191974	402.960219	
min	17.000000	1.228500e+04	1.000000	0.000000	0.000000	
25%	28.000000	1.178270e+05	9.000000	0.000000	0.000000	
50%	37.000000	1.783560e+05	10.000000	0.000000	0.000000	
75%	48.000000	2.370510e+05	12.000000	0.000000	0.000000	
max	90.000000	1.484705e+06	16.000000	14084.000000	4356.000000	

  

	hours-per-week
count	32236.000000
mean	40.450428
std	12.353748
min	1.000000
25%	40.000000
50%	40.000000
75%	45.000000
max	99.000000

```
In [5]: # check missingness, find null values
df.isnull().values.any()
nan_count = np.sum(df.isnull(), axis = 0)
nan_count
```

```
Out[5]: age          162
workclass      1836
fnlwgt         0
education      0
education-num  0
marital-status 0
occupation    1843
relationship   0
race           0
sex_selfID     0
capital-gain   0
capital-loss   0
hours-per-week 325
native-country 583
income_binary  0
dtype: int64
```

```
In [7]: # drop columns that will not be necessary
exclude = ['fnlwgt', 'education-num', 'education', 'marital-status', 'relationship']
newdf = df.drop(columns=exclude, axis=0)

print(newdf)
```

	age	workclass	occupation	race	sex_selfID	\
0	39.0	State-gov	Adm-clerical	White	Non-Female	
1	50.0	Self-emp-not-inc	Exec-managerial	White	Non-Female	
2	38.0	Private	Handlers-cleaners	White	Non-Female	
3	53.0	Private	Handlers-cleaners	Black	Non-Female	
4	28.0	Private	Prof-specialty	Black	Female	
...	...	...	...	...	...	
32556	27.0	Private	Tech-support	White	Female	
32557	40.0	Private	Machine-op-inspct	White	Non-Female	
32558	58.0	Private	Adm-clerical	White	Female	
32559	22.0	Private	Adm-clerical	White	Non-Female	
32560	52.0	Self-emp-inc	Exec-managerial	White	Female	

  

	capital-gain	capital-loss	hours-per-week	native-country	income_binary
0	2174	0	40.0	United-States	<=50K
1	0	0	13.0	United-States	<=50K
2	0	0	40.0	United-States	<=50K
3	0	0	40.0	United-States	<=50K
4	0	0	40.0	Cuba	<=50K
...	...	...	...	...	...
32556	0	0	38.0	United-States	<=50K
32557	0	0	40.0	United-States	>50K
32558	0	0	40.0	United-States	<=50K
32559	0	0	20.0	United-States	<=50K
32560	14084	0	40.0	United-States	>50K

[32561 rows x 10 columns]

```
In [8]: newdf.head()
```

```
Out[8]:
```

	age	workclass	occupation	race	sex_selfID	capital-gain	capital-loss	hours-per-week	native-country	income_binary
0	39.0	State-gov	Adm-clerical	White	Non-Female	2174	0	40.0	United-States	<=50K
1	50.0	Self-emp-not-inc	Exec-managerial	White	Non-Female	0	0	13.0	United-States	<=50K
2	38.0	Private	Handlers-cleaners	White	Non-Female	0	0	40.0	United-States	<=50K
3	53.0	Private	Handlers-cleaners	Black	Non-Female	0	0	40.0	United-States	<=50K
4	28.0	Private	Prof-specialty	Black	Female	0	0	40.0	Cuba	<=50K

## Part 3: Implement Your Project Plan

**Task:** Use the rest of this notebook to carry out your project plan. You will:

1. Prepare your data for your model and create features and a label.
2. Fit your model to the training data and evaluate your model.
3. Improve your model by performing model selection and/or feature selection techniques to find best model for your problem.

Add code cells below and populate the notebook with commentary, code, analyses, results, and figures as you see fit.

## Step 1: Data Preparation and Features:

```
In [11]: # 'income' is the target variable we want to predict

# Separate features (X) and the target variable (y)
X = df.drop('income_binary', axis=1)
y = df['income_binary']

# Perform one-hot encoding on categorical features
X_encoded = pd.get_dummies(X)

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X_encoded, y, test_size=0.2, random_state=42)
```

```
In [13]: # Calculate the mean of each column and store it in a variable
mean_values = df.mean()

# Replace missing values in each column with the corresponding mean value
df.fillna(mean_values, inplace=True)
```

## Step 2: Model Fitting and Evaluation:

```
In [15]: # Initialize and train the Logistic Regression model
```

```
lr = LogisticRegression()
lr.fit(X_train, y_train)

# Predict on the test set
y_pred = lr.predict(X_test)

# Evaluate model performance
accuracy = accuracy_score(y_test, y_pred)
roc_auc = roc_auc_score(y_test, y_pred)
```

```
print("Accuracy:", accuracy)
print("ROC AUC:", roc_auc)
```

```
# Plot ROC curve
plot_roc_curve(lr, X_test, y_test)
plt.title('ROC Curve')
plt.show()
```

```
-----
ValueError                                Traceback (most recent call last)
```

```
<ipython-input-15-3ce0acd39d13> in <module>()
```

```
1 # Initialize and train the Logistic Regression model
2 lr = LogisticRegression()
----> 3 lr.fit(X_train, y_train)
4
5 # Predict on the test set
```

```
/usr/local/lib/python3.6/dist-packages/sklearn/linear_model/_logistic.py in fit(self, X, y, sample_weight)
```

```
1525
1526     X, y = check_X_y(X, y, accept_sparse='csr', dtype=dtype, order="C",
-> 1527                      accept_large_sparse=solver != 'liblinear')
1528     check_classification_targets(y)
1529     self.classes_ = np.unique(y)
```

```
/usr/local/lib/python3.6/dist-packages/sklearn/utils/validation.py in check_X_y(X, y, accept_sparse, accept_large_sparse, dtype, order, copy, force_all_finite, ensure_2d, allow_nd, multi_output, ensure_min_samples, ensure_min_features, y_numeric, warn_on_dtype, estimator)
```

```
753         ensure_min_features=ensure_min_features,
754         warn_on_dtype=warn_on_dtype,
-> 755         estimator=estimator)
756     if multi_output:
757         y = check_array(y, 'csr', force_all_finite=True, ensure_2d=False,
```

```
/usr/local/lib/python3.6/dist-packages/sklearn/utils/validation.py in check_array(array, accept_sparse, accept_large_sparse, dtype, order, copy, force_all_finite, ensure_2d, allow_nd, ensure_min_samples, ensure_min_features, warn_on_dtype, estimator)
```

```
576         if force_all_finite:
577             _assert_all_finite(array,
-> 578                               allow_nan=force_all_finite == 'allow-nan')
579
580     if ensure_min_samples > 0:
```

```
/usr/local/lib/python3.6/dist-packages/sklearn/utils/validation.py in _assert_all_finite(X, allow_nan, msg_dtype)
```

```
58     msg_err.format
59     (type_err,
-> 60     msg_dtype if msg_dtype is not None else X.dtype)
61 )
62 # for object dtype data, we only check for NaNs (GH-13254)
```

```
ValueError: Input contains NaN, infinity or a value too large for dtype('float64').
```

```
In [16]: param_grid = {
        'C': [0.001, 0.01, 0.1, 1, 10],
        'penalty': ['l1', 'l2']
    }

    grid_search = GridSearchCV(lr, param_grid, cv=5, scoring='roc_auc')
    grid_search.fit(X_train, y_train)

    # Get the best model and its parameters
    best_logreg = grid_search.best_estimator_
    best_params = grid_search.best_params_

    # Re-train the best model on the entire training set
    best_logreg.fit(X_train, y_train)

    # Evaluate the improved model
    y_pred_best = best_logreg.predict(X_test)
    accuracy_best = accuracy_score(y_test, y_pred_best)
    roc_auc_best = roc_auc_score(y_test, y_pred_best)

    print("Best Model Accuracy:", accuracy_best)
    print("Best Model ROC AUC:", roc_auc_best)

    # Plot ROC curve for the best model
    plot_roc_curve(best_logreg, X_test, y_test)
    plt.title('ROC Curve (Best Model)')
    plt.show()
```



```
/usr/local/lib/python3.6/dist-packages/sklearn/model_selection/_validation.py:536: FitFailedWarning: Estimator fit failed. The score on this train-test partition for these parameters will be set to nan. Details:
ValueError: Input contains NaN, infinity or a value too large for dtype('float64').
```

```
FitFailedWarning)
/usr/local/lib/python3.6/dist-packages/sklearn/model_selection/_validation.py:536: FitFailedWarning: Estimator fit failed. The score on this train-test partition for these parameters will be set to nan. Details:
ValueError: Input contains NaN, infinity or a value too large for dtype('float64').
```

```
FitFailedWarning)
/usr/local/lib/python3.6/dist-packages/sklearn/model_selection/_validation.py:536: FitFailedWarning: Estimator fit failed. The score on this train-test partition for these parameters will be set to nan. Details:
ValueError: Input contains NaN, infinity or a value too large for dtype('float64').
```

```
FitFailedWarning)
/usr/local/lib/python3.6/dist-packages/sklearn/model_selection/_validation.py:536: FitFailedWarning: Estimator fit failed. The score on this train-test partition for these parameters will be set to nan. Details:
ValueError: Input contains NaN, infinity or a value too large for dtype('float64').
```

```
FitFailedWarning)
/usr/local/lib/python3.6/dist-packages/sklearn/model_selection/_validation.py:536: FitFailedWarning: Estimator fit failed. The score on this train-test partition for these parameters will be set to nan. Details:
ValueError: Input contains NaN, infinity or a value too large for dtype('float64').
```

```
FitFailedWarning)
/usr/local/lib/python3.6/dist-packages/sklearn/model_selection/_validation.py:536: FitFailedWarning: Estimator fit failed. The score on this train-test partition for these parameters will be set to nan. Details:
ValueError: Solver lbfgs supports only 'l2' or 'none' penalties, got l1 penalty.
```

```
FitFailedWarning)
/usr/local/lib/python3.6/dist-packages/sklearn/model_selection/_validation.py:536: FitFailedWarning: Estimator fit failed. The score on this train-test partition for these parameters will be set to nan. Details:
ValueError: Input contains NaN, infinity or a value too large for dtype('float64').
```

```
FitFailedWarning)
/usr/local/lib/python3.6/dist-packages/sklearn/model_selection/_validation.py:536: FitFailedWarning: Estimator fit failed. The score on this train-test partition for these parameters will be set to nan. Details:
ValueError: Input contains NaN, infinity or a value too large for dtype('float64').
```

```
FitFailedWarning)
/usr/local/lib/python3.6/dist-packages/sklearn/model_selection/_validation.py:536: FitFailedWarning: Estimator fit failed. The score on this train-test partition for these parameters will be set to nan. Details:
ValueError: Input contains NaN, infinity or a value too large for dtype('float64').
```

```
FitFailedWarning)
/usr/local/lib/python3.6/dist-packages/sklearn/model_selection/_validation.py:536: FitFailedWarning: Estimator fit failed. The score on this train-test partition for these parameters will be set to nan. Details:
ValueError: Input contains NaN, infinity or a value too large for dtype('float64').
```

```
FitFailedWarning)
/usr/local/lib/python3.6/dist-packages/sklearn/model_selection/_validation.py:536: FitFailedWarning: Estimator fit failed. The score on this train-test partition for these parameters will be set to nan. Details:
ValueError: Input contains NaN, infinity or a value too large for dtype('float64').
```

```
FitFailedWarning)
```

```
-----
ValueError                                Traceback (most recent call last)
<ipython-input-16-53bc0792f786> in <module>()
      5
      6 grid_search = GridSearchCV(lr, param_grid, cv=5, scoring='roc_auc')
----> 7 grid_search.fit(X_train, y_train)
      8
      9 # Get the best model and its parameters

/usr/local/lib/python3.6/dist-packages/sklearn/model_selection/_search.py in fit(self, X, y, groups, **fit_params)
    737         refit_start_time = time.time()
    738         if y is not None:
--> 739             self.best_estimator_.fit(X, y, **fit_params)
    740         else:
    741             self.best_estimator_.fit(X, **fit_params)

/usr/local/lib/python3.6/dist-packages/sklearn/linear_model/_logistic.py in fit(self, X, y, sample_weight)
    1486         The SAGA solver supports both float64 and float32 bit arrays.
    1487         """
-> 1488         solver = _check_solver(self.solver, self.penalty, self.dual)
    1489
    1490         if not isinstance(self.C, numbers.Number) or self.C < 0:

/usr/local/lib/python3.6/dist-packages/sklearn/linear_model/_logistic.py in _check_solver(solver, penalty, dual)
    443         if solver not in ['liblinear', 'saga'] and penalty not in ('l2', 'none'):
    444             raise ValueError("Solver %s supports only 'l2' or 'none' penalties, "
--> 445                               "got %s penalty." % (solver, penalty))
    446         if solver != 'liblinear' and dual:
    447             raise ValueError("Solver %s supports only "

ValueError: Solver lbfgs supports only 'l2' or 'none' penalties, got l1 penalty.
```

```
In [ ]:
```